Rain prediction using ANN

Ahmed Alelg

Problem statement

To predicate whether it will rain or not in Australia.

Dataset

- Outsourced from Kaggle.
- 23 Columns
  - 1 Datetime
  - 15 Numerical
  - 7 Categorical
- 145,460 observations

```
 #   Column         Non-Null Count   Dtype
---  ------         --------------   -----
 0   Date           145460 non-null  datetime64[ns]
 1   Location       145460 non-null  object
 2   MinTemp        143975 non-null  float64
 3   MaxTemp        144199 non-null  float64
 4   Rainfall       142199 non-null  float64
 5   Evaporation    82670 non-null   float64
 6   Sunshine       75625 non-null   float64
 7   WindGustDir    135134 non-null  object
 8   WindGustSpeed  135197 non-null  float64
 9   WindDir9am     134894 non-null  object
 10  WindDir3pm     141232 non-null  object
 11  WindSpeed9am   143693 non-null  float64
 12  WindSpeed3pm   142398 non-null  float64
 13  Humidity9am    142806 non-null  float64
 14  Humidity3pm    140953 non-null  float64
 15  Pressure9am    130395 non-null  float64
 16  Pressure3pm    130432 non-null  float64
 17  Cloud9am       89572 non-null   float64
 18  Cloud3pm       86102 non-null   float64
 19  Temp9am        143693 non-null  float64
 20  Temp3pm        141851 non-null  float64
 21  RainToday      142199 non-null  object
 22  RainTomorrow   142193 non-null  object
dtypes: datetime64[ns](1), float64(16), object(6)
memory usage: 25.5+ MB
```

EDA

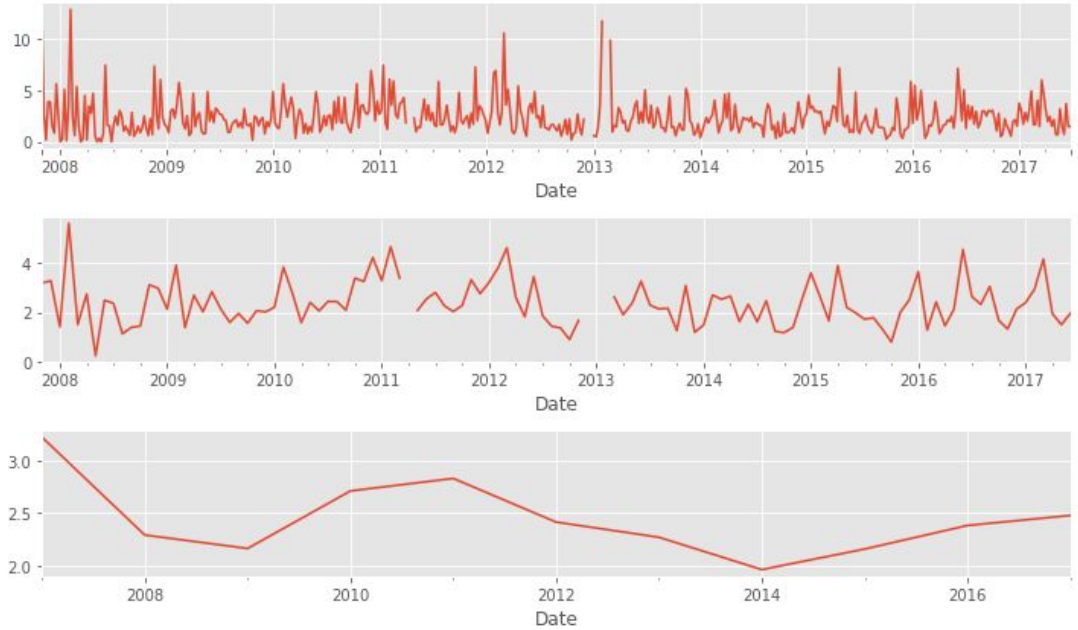
10
5
0
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
Date
4
2
0
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
Date
3.0
2.5
2.0
2008
2010
2012
2014
2016
Date

EDA

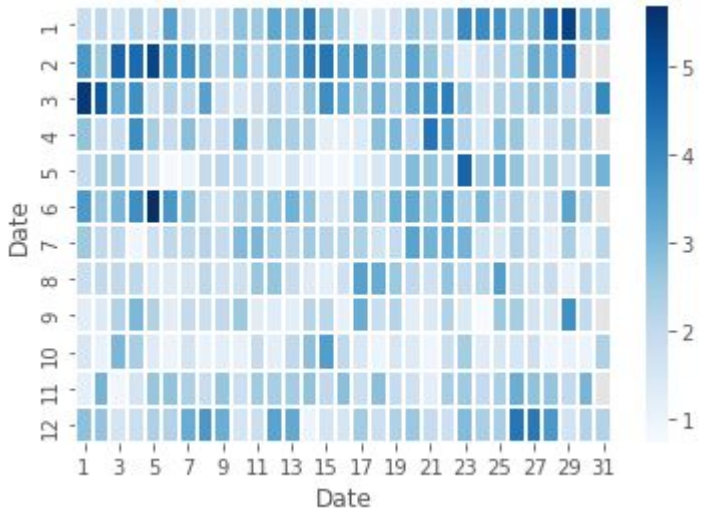
Date
1
2
3
4
5
6
7
8
9
10
11
12
1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 31
Date
5
4
3
2
1

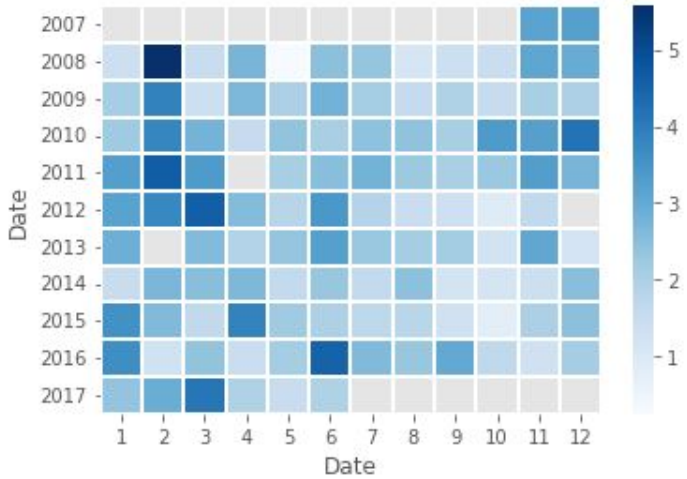
Date
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
1 2 3 4 5 6 7 8 9 10 11 12
Date
5
4
3
2
1

EDA

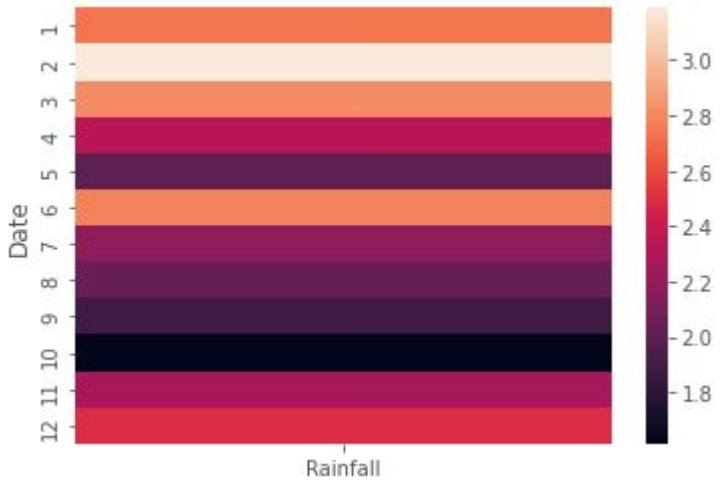
Date
1
2
3
4
5
6
7
8
9
10
11
12
Rainfall
3.0
2.8
2.6
2.4
2.2
2.0
1.8

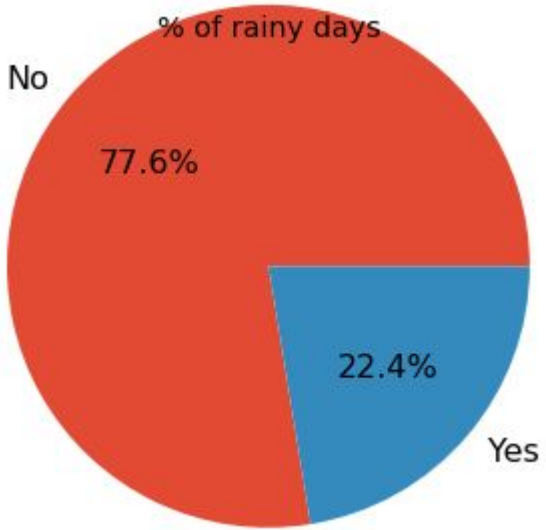
% of rainy days
No
77.6%
22.4%
Yes

EDA

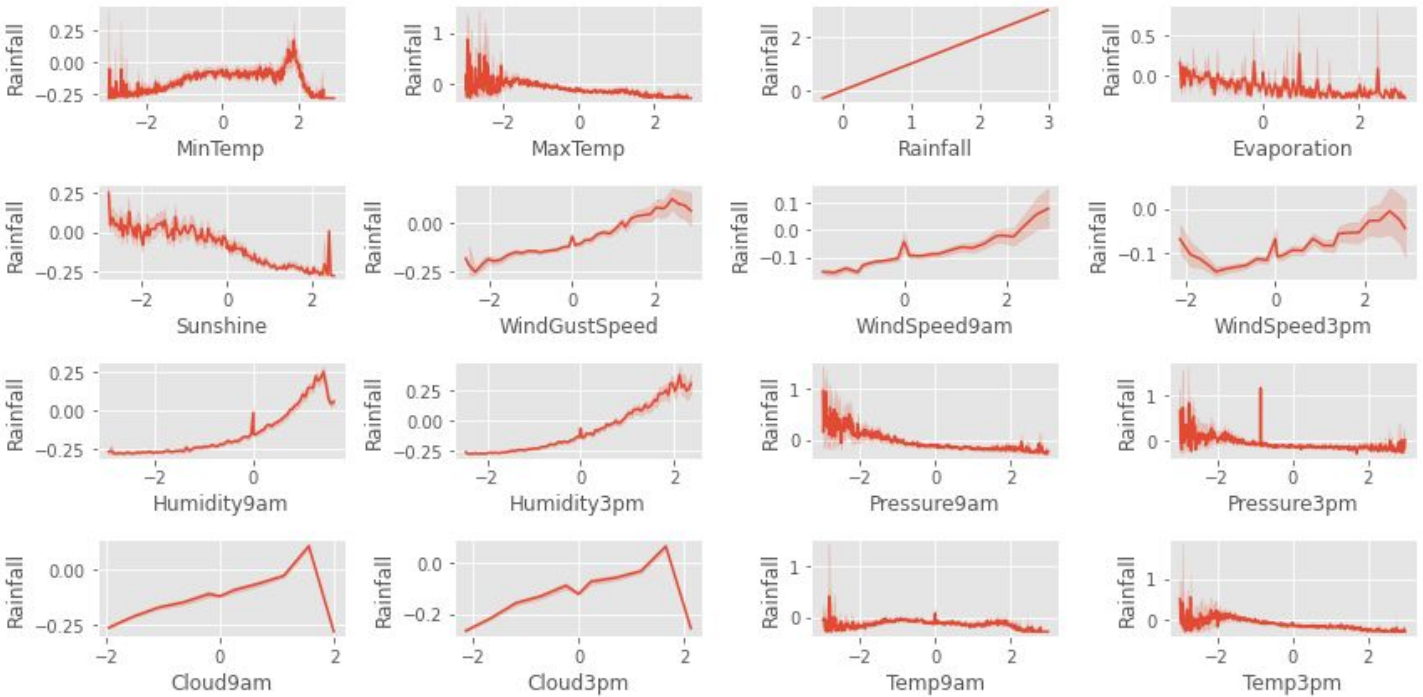
Rainfall
MinTemp
MaxTemp
Rainfall
Evaporation
Sunshine
WindGustSpeed
WindSpeed9am
WindSpeed3pm
Humidity9am
Humidity3pm
Pressure9am
Pressure3pm
Cloud9am
Cloud3pm
Temp9am
Temp3pm

EDA

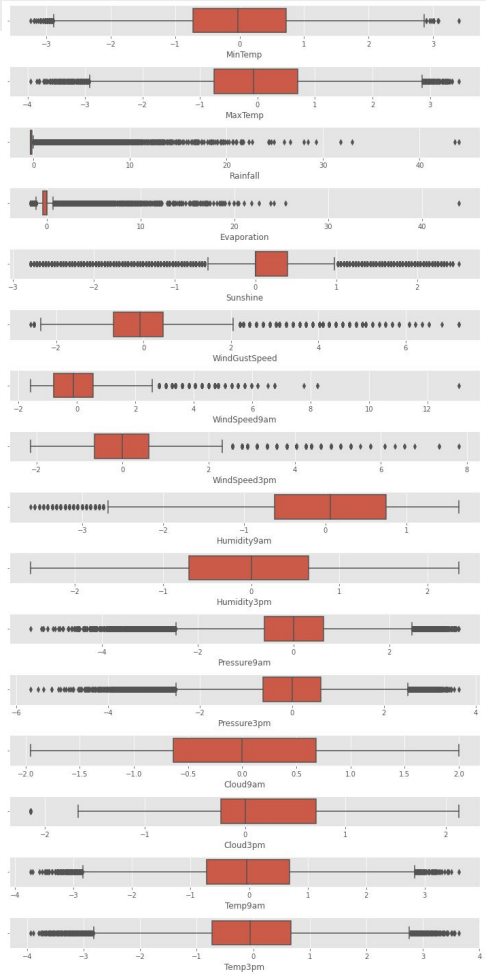
MinTemp
MaxTemp
Rainfall
Evaporation
Sunshine
WindGustSpeed
WindSpeed9am
WindSpeed3pm
Humidity9am
Humidity3pm
Pressure9am
Pressure3pm
Cloud9am
Cloud3pm
Temp9am
Temp3pm

Data engineering

- Many columns with missing values!
  - Numerical values filled with the mean
  - Categorical values filled with the its neighbor.
  - Columns having >50% missing values were dropped.

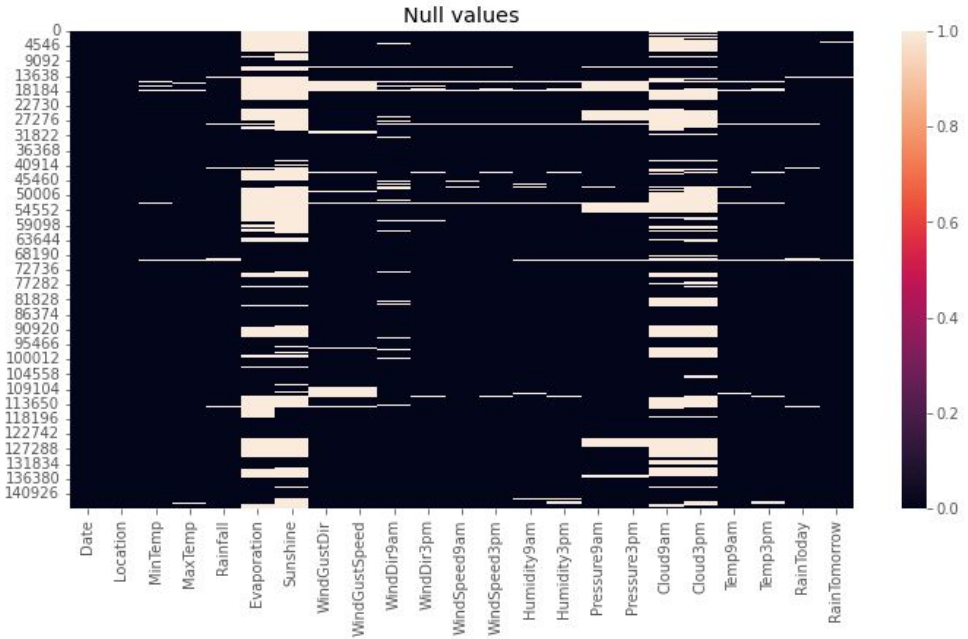

Data engineering

- Standardized the numerical values to have range within -1,1.
- Removed outliers whose values exceeded [-3, 3]
  - n = 8,674
- Removed columns with high collinearity VIF > 5.
  - n = 4

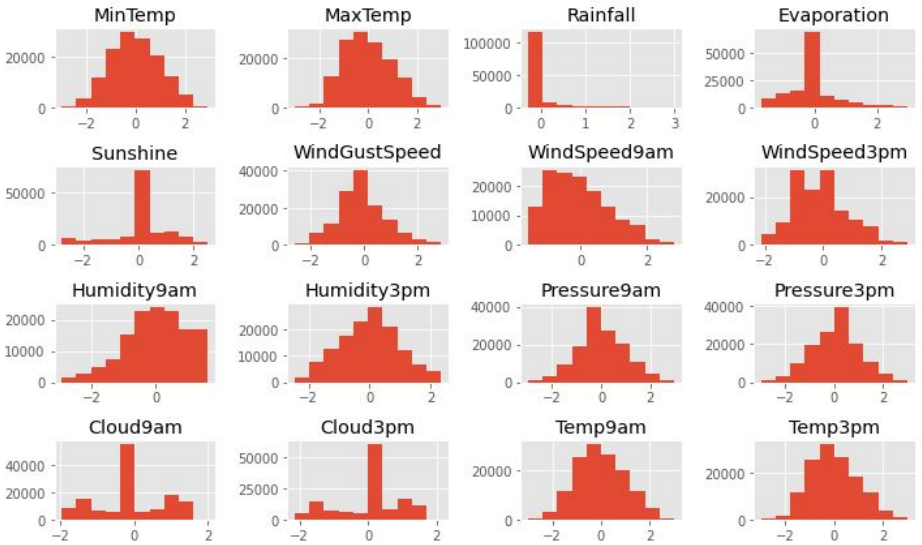

Data engineering

- Encoded categorical values to integers.
- Created vector embeddings for categorical values.
- Encoded Month/Day values into cyclical features.

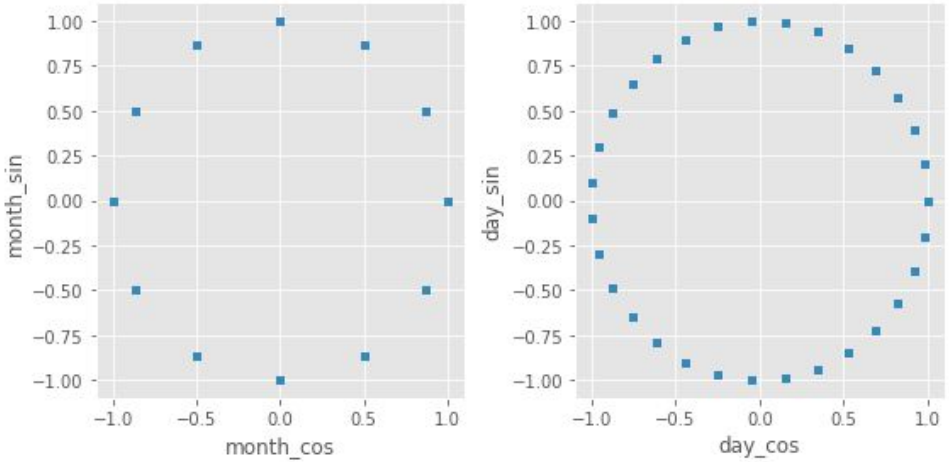

Data split & balance

- 90% Training data
  - Balanced classes using SMOTE
  - (n = 189,680)
- 5% Validation
- 5% Testing

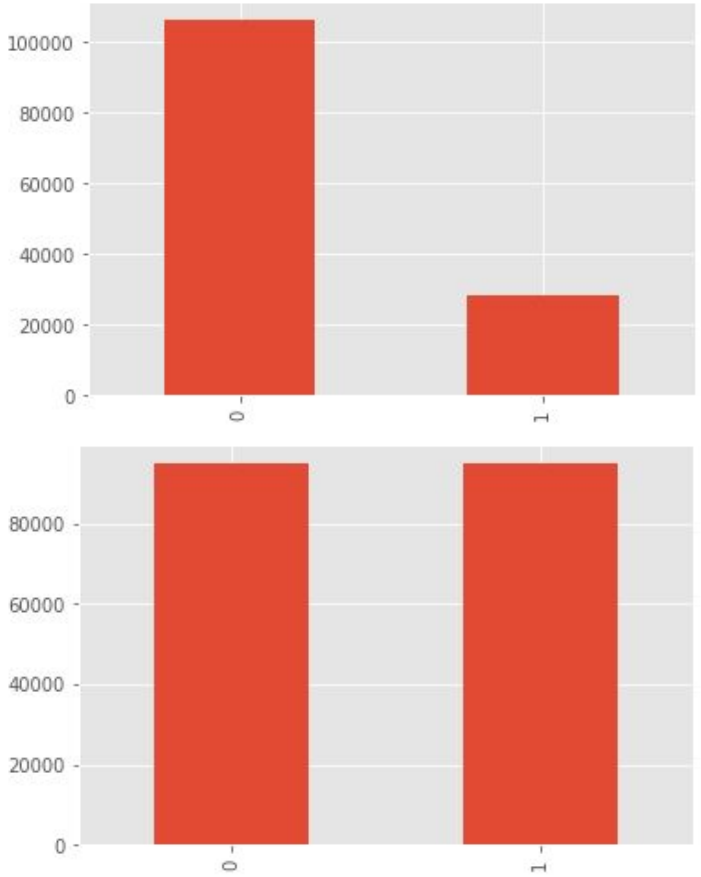

Model architecture

- ANN with 6 layers
- Used embeddings to encode categorical data.
- ReLU activation functions.
- Included dropout and Batch normalization.
- Adam.
- Binary cross entropy loss.

```
NNet(
  (embds): ModuleList(
    (0): Embedding(49, 25)
    (1): Embedding(16, 8)
    (2): Embedding(16, 8)
    (3): Embedding(16, 8)
  )
  (emb_drop): Dropout(p=0.05, inplace=False)
  (act): ReLU()
  (init): Sequential(
    (0): Linear(in_features=69, out_features=138, bia
    (1): BatchNorm1d(138, eps=1e-05, momentum=0.1, af
    (2): ReLU()
    (3): Dropout(p=0.3, inplace=False)
    (4): Linear(in_features=138, out_features=256, bi
    (5): BatchNorm1d(256, eps=1e-05, momentum=0.1, af
    (6): ReLU()
    (7): Dropout(p=0.3, inplace=False)
    (8): Linear(in_features=256, out_features=128, bi
    (9): BatchNorm1d(128, eps=1e-05, momentum=0.1, af
    (10): ReLU()
    (11): Dropout(p=0.3, inplace=False)
    (12): Linear(in_features=128, out_features=64, bi
    (13): BatchNorm1d(64, eps=1e-05, momentum=0.1, af
    (14): ReLU()
    (15): Dropout(p=0.3, inplace=False)
    (16): Linear(in_features=64, out_features=16, bia
    (17): BatchNorm1d(16, eps=1e-05, momentum=0.1, af
    (18): ReLU()
    (19): Dropout(p=0.1, inplace=False)
```

Results

	Training	Validation	Testing
Loss	0.339279	0.365065	0.362375
Accuracy	84.71%	84.23%	84.02%
F1	0.8428	0.6689	0.6517

Results: Test data

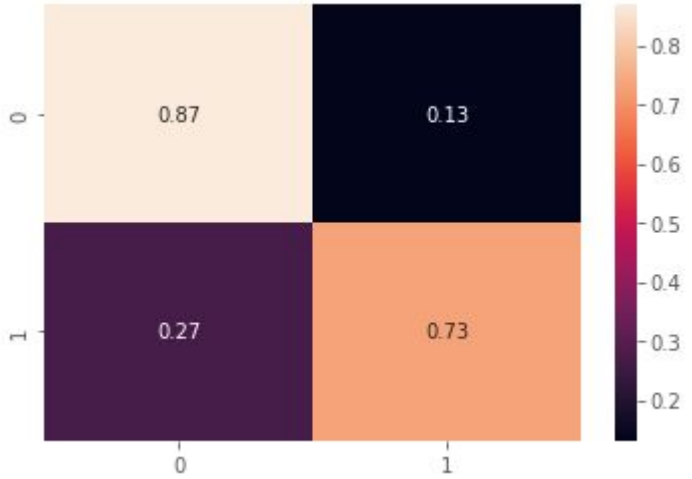

Results: Random Forest Classifier

	Training	Validation	Testing
Accuracy	100%	84.89%	85.58%
F1	1.0	0.6473	0.6448

Results: Test data

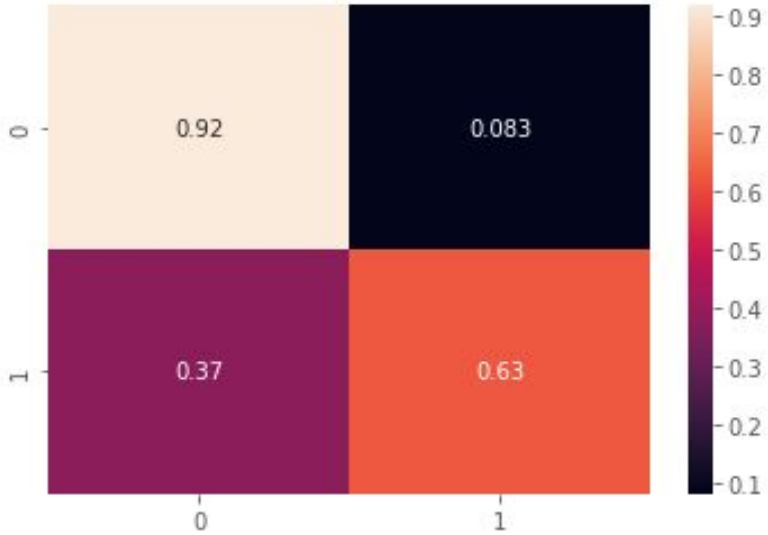

Libraries used

- Numpy
- Pandas
- Sklearn
- PyTorch
- imbalanced-learn
- seaborn & matplotlib